

C++ Map-like Data Structures: Types, Pros/Cons, and Performance

std::map

Keywords: Balanced Binary Search Tree (Red-Black Tree), ordered, unique keys

Pros:

- Keys kept in sorted order
- Efficient for ordered traversal
- Deterministic performance ($O(\log n)$)

Cons:

- Slower than hash maps for random lookups
- Higher memory usage than flat structures

Time Complexity:

Operation	Complexity
Lookup	$O(\log n)$
Insertion	$O(\log n)$
Deletion	$O(\log n)$
Iteration	$O(n)$

std::multimap

Keywords: Balanced Binary Search Tree, ordered, allows duplicate keys

Pros:

- Keys stored in sorted order
- Supports multiple values per key
- Good for grouped data

Cons:

- Same drawbacks as std::map (slower lookups than hash maps)
- Duplicates may add overhead

Time Complexity:

Operation	Complexity
Lookup	$O(\log n)$
Insertion	$O(\log n)$
Deletion	$O(\log n)$
Iteration	$O(n)$

std::unordered_map

Keywords: Hash table, unique keys, not ordered

Pros:

- Average constant-time lookup ($O(1)$)
- Faster than std::map for direct key access
- Good for large datasets

Cons:

- No ordering of keys
- Worst-case $O(n)$ if hash collisions occur
- More memory overhead due to buckets

Time Complexity:

Operation	Complexity
Lookup	Average $O(1)$, Worst $O(n)$
Insertion	Average $O(1)$, Worst $O(n)$
Deletion	Average $O(1)$, Worst $O(n)$
Iteration	$O(n)$

std::unordered_multimap

Keywords: Hash table, allows duplicate keys, not ordered

Pros:

- Allows multiple values per key
- Fast average lookups ($O(1)$)
- Good when duplicates are needed

Cons:

- No key ordering

- More memory overhead
- Worst-case $O(n)$ with many collisions

Time Complexity:

Operation	Complexity
Lookup	Average $O(1)$, Worst $O(n)$
Insertion	Average $O(1)$, Worst $O(n)$
Deletion	Average $O(1)$, Worst $O(n)$
Iteration	$O(n)$

boost::unordered_flat_map

Keywords: Hybrid: Hash table + flat contiguous storage

Pros:

- Fast lookups (average $O(1)$)
- Better cache locality than `std::unordered_map`
- Lower memory overhead

Cons:

- Insertions may be slower due to relocations
- Requires Boost library

Time Complexity:

Operation	Complexity
Lookup	Average $O(1)$
Insertion	Average $O(1)$
Deletion	Average $O(1)$
Iteration	$O(n)$, faster than <code>std::unordered_map</code>

boost::flat_map

Keywords: Sorted vector-like container, ordered keys, unique

Pros:

- Very memory efficient
- Fast iteration due to contiguous storage
- Good for small to medium datasets

Cons:

- Insertions/deletions are $O(n)$ (need to shift elements)
- Not suitable for very large or frequently updated datasets

Time Complexity:

Operation	Complexity
Lookup	$O(\log n)$
Insertion	$O(n)$
Deletion	$O(n)$
Iteration	$O(n)$, but very fast due to cache locality