

# Dev Cheat Sheet — `iovec`, `writew/pwritev`, `off_t`, `size_t`, `uint64_t`, and related C++ keywords

A practical, zero-fluff guide so you know **what each thing is**, **when to use it**, **why**, and the **pros/cons**—with simple, practical examples.

---

## 1) POSIX scatter/gather I/O: `iovec`, `writew`, `pwritew`

What is `iovec`?

- A tiny struct ( `<sys/uio.h>` ) that **points to a buffer** in memory: base pointer + length.
- You build an **array of** `iovec` values and pass them to `writew` / `pwritew` to write **multiple buffers in one syscall**.

```
struct iovec { void* iov_base; size_t iov_len; };
```

`writew` vs `pwritew`

- `writew(fd, iov, iovcnt)` writes many buffers sequentially **at the current file offset**.
- `pwritew(fd, iov, iovcnt, offset)` writes many buffers **at an explicit file offset** (does not change the file pointer; thread-safe for shared fds).

### When to use

- You have **many small pieces** (e.g., header + payload, or many entries) and want **one syscall** instead of many.
- In practice: batch multiple `{header, data}` pairs for ops/index flushes to minimize open/close/seek/write overhead.

### Pros

- Fewer syscalls; better throughput and CPU efficiency.
- Kernel can optimize the combined write.

### Cons / gotchas

- You must ensure each `iov_base` points to **valid, still-alive memory** until the call returns.
  - Very large `iovcnt` may hit OS limits (e.g., `IOV_MAX`). Batch sensibly.
-

## 2) File positions & sizes: `off_t` vs `size_t`

### `off_t` (file offset type)

- Signed integer type used by POSIX to represent **byte offsets** in a file.
- Why signed? Historically ties to `lseek` semantics; modern use is just “position in the file”.
- Use `off_t` for: where to read/write (e.g., `ops_file_position_`, `index_file_position_`).

**Pros:** Standard for file offsets; large-file ready on 64-bit builds. **Cons:** Signed; be careful with negative checks and conversions.

### `size_t` (object size/length type)

- Unsigned type returned by `sizeof` and container `.size()`. Represents **sizes and counts**.
- Use `size_t` for: how many bytes (e.g., `ops_file_reserved_size_`, buffer lengths, vector sizes).

**Pros:** Natural result of `sizeof` and STL size APIs; wide on 64-bit. **Cons:** Unsigned—be careful with underflow and signed/unsigned comparisons.

**Rule of thumb:** - **Where?** Use `off_t` for positions in a file. - **How much?** Use `size_t` for lengths/sizes.

---

## 3) Big integer choices: `uint64_t`, `int64_t`, `uint32_t`, `int`

### `uint64_t`

- Unsigned 64-bit integer (0 ...  $2^{64}-1$ ). Great for **IDs, counters, sizes** that should never be negative.
- In practice: `file_id_` is `uint64_t` so we can scale to **huge numbers of files/ops** without overflow.

### Peers & when to use them

- `int64_t` → need negatives and large range (e.g., deltas, signed math over large domain).
- `uint32_t` / `int32_t` → smaller ranges; good for compact structures or network formats.
- `int` → convenient but **width varies** by platform/ABI; avoid for on-disk/network formats.

**Pros of fixed-width types** (`uint64_t`, etc.): deterministic size across platforms. **Cons:** Slightly more verbose; may use more memory than necessary if over-provisioned.

---

## 4) C++ keywords you see in these structs

final

- `struct X final { ... };` → **prevents inheritance** from `X`.
- Rationale: freeze the data model; avoid vtables/ABI surprises for simple POD-like structs.

static

- Member belongs to the **type**, not each instance.

constexpr

- Value is a **compile-time constant**. Enables optimizations and compile-time checks.

auto

- **Type deduction** by the compiler. Here used with `constexpr` for a named flag.

### Example in your code

```
struct ops_index_flush_entry final {  
    static constexpr auto Reflect = true; // compile-time flag for reflection  
    utilities  
    ...  
};
```

**Why this exists:** your logging/serialization helpers (e.g., `reflect::toString`) can gate behavior on `Reflect`.

## 5) Two example “flush entry” structs & their roles

OpsFlushEntry

- **What:** one ops update: which file, which ops file, **where** to write it (`off_t`), **how big** (`size_t`), and the **data buffer**.
- **Used by:** a function that writes ops entries to disk.
- **Goal:** build all `{header, data}` iovecs **first**, then do **batched** `pwritev` per target file via your I/O layer.

IndexFlushEntry

- **What:** the index-level record that points to the ops block (filename + offset + reserved size) for a `file_id`.
- **Used by:** a function that writes index entries to an index file.

- **Goal:** similarly batch writes to the index file.

---

## 6) How to choose the right type quickly (decision guide)

- Need to mark a byte **POSITION** in a file? → `off_t`.
- Need a **BYTES LENGTH** or container size? → `size_t`.
- Need a huge non-negative ID/counter (won't be negative)? → `uint64_t`.
- Need large numbers that may be negative? → `int64_t`.
- Writing many buffers at once to a file? → Build `std::vector<iovec>` and call `pwritev`.
- Writing at a specific offset without moving the file pointer? → `pwritev` (not `writev`).

---

## 7) Mini examples

### 7.1 Build iovecs and use `pwritev`

```
#include <sys/uio.h>
#include <unistd.h>

ssize_t write_header_and_data(int fd, off_t at, const std::string& hdr, const
std::string& data) {
    iovec vecs[2] = {
        { const_cast<char*>(hdr.data()), hdr.size() },
        { const_cast<char*>(data.data()), data.size() }
    };
    return ::pwritev(fd, vecs, 2, at);
}
```

### 7.2 `off_t` vs `size_t`

```
off_t position = 4096;    // start 4 KiB into the file
size_t length  = 1024;    // write 1 KiB
```

### 7.3 Big IDs

```
uint64_t file_id = 9876543210123456789ULL; // plenty of room for growth
```

## 8) Pros & cons summary table

Concept	Use for	Pros	Cons
<code>iovec</code> + <code>pwritev</code>	One syscall for many buffers	Fewer syscalls, better throughput	Must manage lifetimes; <code>IOV_MAX</code> cap
<code>off_t</code>	File offsets	Large-file ready, standard	Signed; watch negative values
<code>size_t</code>	Sizes/lengths	Matches STL APIs, wide on 64-bit	Unsigned pitfalls in comparisons
<code>uint64_t</code>	IDs/counters (non-negative)	Huge range, portable width	More memory than smaller ints
<code>final</code>	Seal a type	Prevents unsafe subclassing	Less flexible for extension
<code>static constexpr</code>	Compile-time flags/constants	Zero run-time cost	Requires compile-time known value

## 9) Applying this to a typical flush pipeline

- **Today:** per entry → open, `pwritev(header, data)`, close.
- **Target:** group entries by target file, prebuild all `iovec`s, **one** `pwritev` **per file** through your I/O layer.
- **Benefits:** significant reduction in open/close/syscall churn; better sequential I/O; cleaner integration with your async I/O pipeline.

## 10) Quick checklist when implementing batching

- [] Group `ops_index_flush_entry` by `ops_file_name_`.
- [] For each group, sort by `ops_file_position_`.
- [] Build a single `std::vector<iovec>` as `{header1, data1, header2, data2, ...}`.
- [] Ensure buffers stay alive until `pwritev` completes.
- [] Respect `IOV_MAX` → split into multiple calls if needed.
- [] Use `pwritev(fd, iovecs.data(), iovecs.size(), start_offset_of_first)` **or** post via `drive_io` op.
- [] Handle partial writes / errors robustly; log and disable bucket on corruption.

# Diagrams

## A) Scatter/Gather I/O with `iovec`

User buffers in RAM

```
[ header1 ] [ data1 ] [ header2 ] [ data2 ] ...  
    |         |         |         |  
    v         v         v         v  
  iovec[0]   iovec[1]   iovec[2]   iovec[3]   ... --> single pwritev(fd,  
iovec[], n, offset)  
  
                                     (kernel writes header1,  
then data1, then header2, ...)
```

**Key idea:** many small pieces → one syscall.

## B) Offsets vs Sizes

```
File (bytes)  
0          4096          5120  
|-----|=====|----->  
^         ^  
|         |  
|         +-- size_t length (e.g., 1024 bytes)  
+-----+ off_t offset (e.g., 4096)
```

- `off_t` = **where** in the file.
- `size_t` = **how much** to read/write.

## C) Batching entries per file

Entries (unsorted):

```
(file A, off 0,    H1,D1)  
(file A, off 4096, H2,D2)  
(file B, off 0,    H3,D3)
```

Group by file → sort by offset → build iovecs per file

```
File A iovecs: [H1][D1][H2][D2] --> pwritev(fd_A, vecA, nA, 0)
File B iovecs: [H3][D3]          --> pwritev(fd_B, vecB, nB, 0)
```

Benefits: fewer opens/closes, fewer syscalls, better locality.

---

## D) Type picking quick guide

Need a file POSITION (byte offset)?

↳ `off_t`

Need a BYTES LENGTH or container size?

↳ `size_t`

Need a huge non-negative ID/counter?

↳ `uint64_t`

↳ If negatives possible → `int64_t`

Writing multiple buffers at once?

↳ `iovec[]` + `writev()`

---

## E) Error handling flow (batched write)

Build iovecs → call `writev` → check return

└ `-1` → log `errno`, retry/backoff or fail

└ `< sum(iov_len)` → partial write: adjust iovecs,

retry remaining